

AHAS : Adaptive Hybrid Accelerator for ScRNA

Shivaritha Sakthi Rengasamy

Date: 12/08/2025





shutterstock.com · 249280997

JEANS

GENES

↖ The Biological Context

Genomic Analysis

Decoding the blueprint of life to understand cellular identity.

Biological Fidelity

Identifying the **rare biological components** amidst the noise.

Identifying Mutations (Signal vs Noise)

A T C G A T C G A T C G A T C G
A T C G A T C G **A** A T C G A T C G
C G A T C G A T C G A T C G

● Housekeeping (Noise) ● **Rare Component**

What Limits Genome Sequence Analysis?

- Significant data movement between storage, CPU, and accelerators.
- Every raw data transfer can amplify the volume of intermediate results.
- The database accumulates noisy, low-value data and grows much faster than the actual compute needs.



PROPOSED SOLUTION

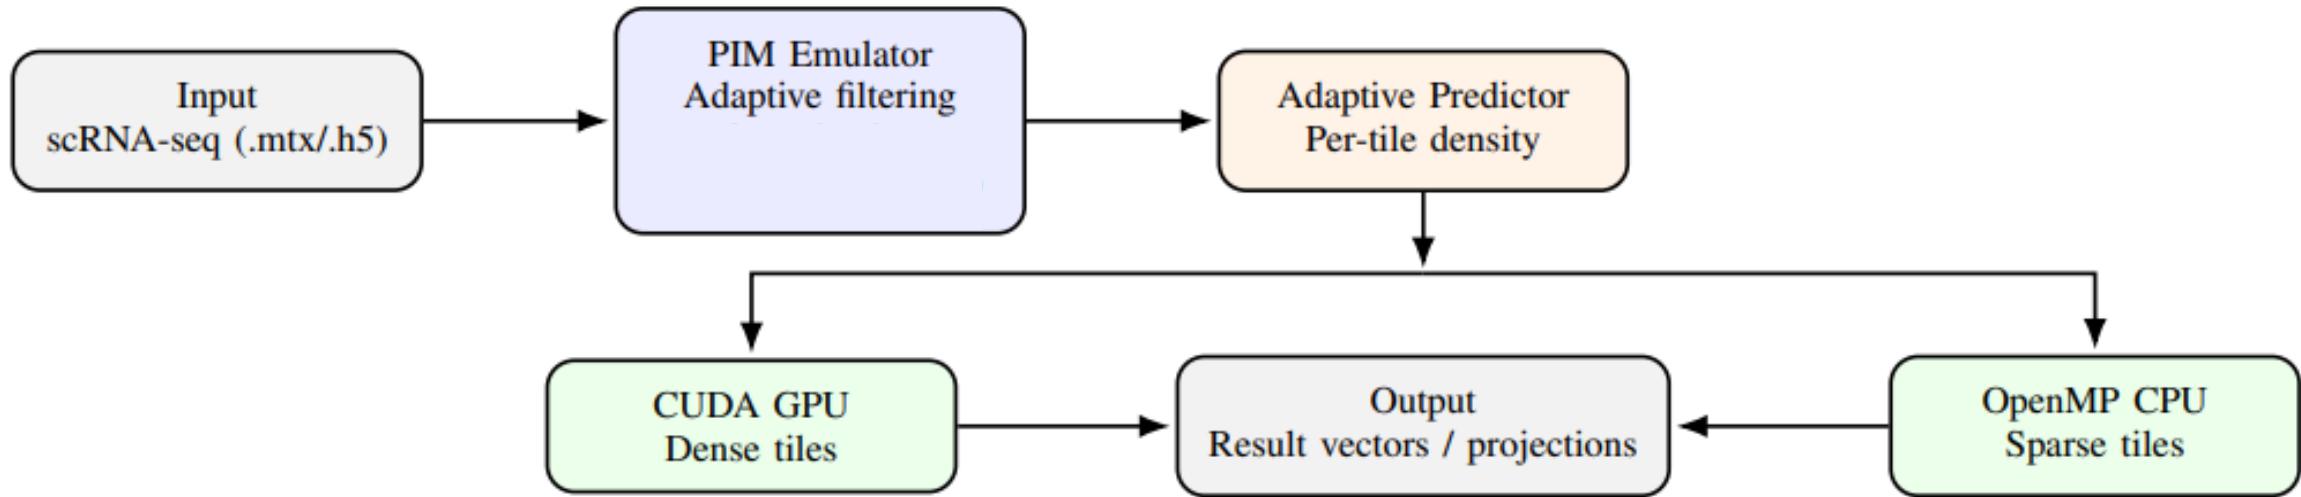


Fig. 1. Hybrid Adaptive SpMM framework.

PIM FILTER



- Variance computation: computes variance per gene across all cells using `np.var` (two-pass: mean, then variance).
- Gene classification: identifies rare genes (top 10% variance), noise genes (bottom 10%), and housekeeping genes (high mean, low variance).
- Filtering: keeps only rare genes (highly variable genes) and saves them to a filtered H5 file.

PIM filtering cuts nonzeros by **about 40-85%** and rows by **around 90%**, so each dataset becomes much smaller before compute.

▀ The Bottleneck: Data Movement

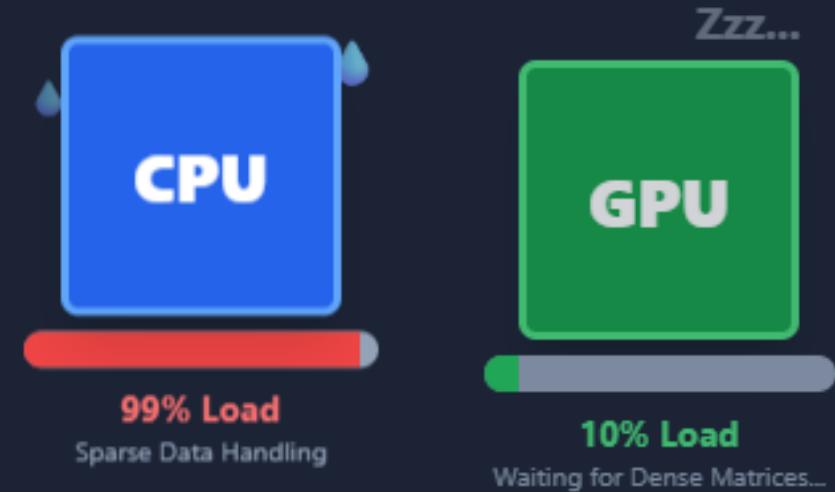
TARGET

Reduce I/O transfers between storage and CPU/GPU.

CHALLENGE

Retain biologically important genes.

Wait... that means the dataset is SPARSE,
not dense?



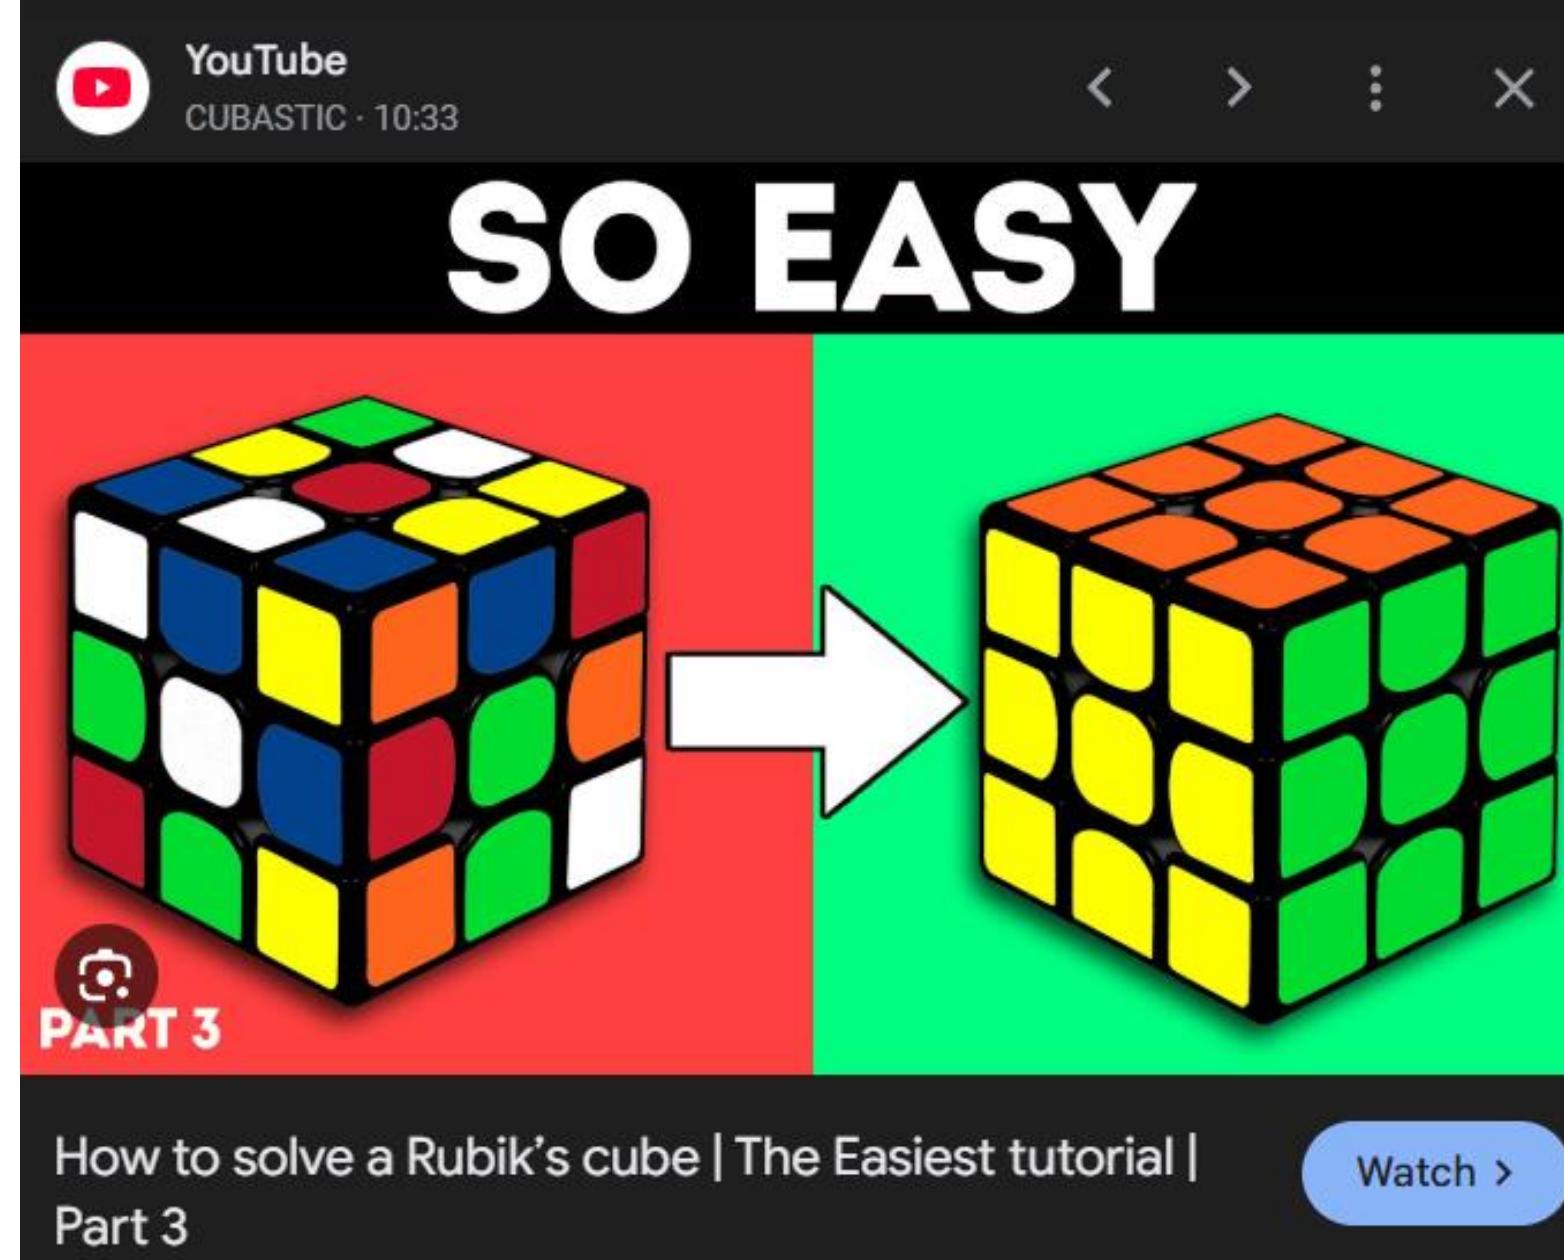
ARE TILES GPU FRIENDLY?

Dense tiles are permuted by sorting rows and columns in descending order by their non-zero counts

The code does this
in `dense_perm_spmm_tile()`:

- Rows are sorted by `nnz_per_row` in descending order
- Columns are sorted by `nnz_per_col` in descending order
- The corresponding W weight matrix rows are permuted to match

This groups dense regions together so the GPU can access memory more efficiently.



Observation

- A sanity check is done between CUDA output vs reference to catch errors.
- IEEE 754: Floating-point addition is not associative; order changes results.
- Example: $(a+b)+c$ can differ from $a+(b+c)$ due to rounding.
- Parallel impact: GPU/parallel code often sums in different orders than CPU.
- Expected drift: Small numeric deltas are normal; large ones indicate issues.
- Tolerances used: $\text{rtol} = 1e-3$ (0.1% relative), $\text{atol} = 1e-5$ (absolute).
- Match rule: Values match if $|a-b| \leq \text{atol} + \text{rtol} * \max(|a|, |b|)$.
- Goal: Ignore harmless floating-point noise; flag true computation errors.

DATASET

Dataset	Dataset Name	Dimensions (rows × cols)	NNZ	Total Elements	Density	Sparsity
1	10k Mouse Splenocytes (5p, GEM-X)	33,696 × 9,263	26,407,826	312,123,648	8.46%	91.54%
2	10k Human DTC Melanoma (NextGEM 5p)	38,606 × 6,704	14,971,913	258,814,624	5.78%	94.22%
3	PBMC from Healthy Donor, Granulocytes Removed (~3k cells)	134,920 × 2,711	24,511,186	365,978,120	6.70%	93.30%
4	30k A549 Lung Carcinoma Cells, CRISPR Pool, Multiplexed	36,706 × 2,979	16,028,697	109,428,174	14.64%	85.36%

DATASET 3 - 2.1x Faster

Metric	Baseline	tilepredpermfspmm	Improvement
Compute time	137.9 ms	66.1 ms	2.1x faster
Input nnz	24.5 M	3.4 M	86% reduction
Input rows	134,920	3,660	97% reduction
Tiles	-	2,494	-
Dense tiles	-	43 (1.7%)	-
Sparse tiles	-	2,451 (98.3%)	-

DATASET 4 - GPU overload

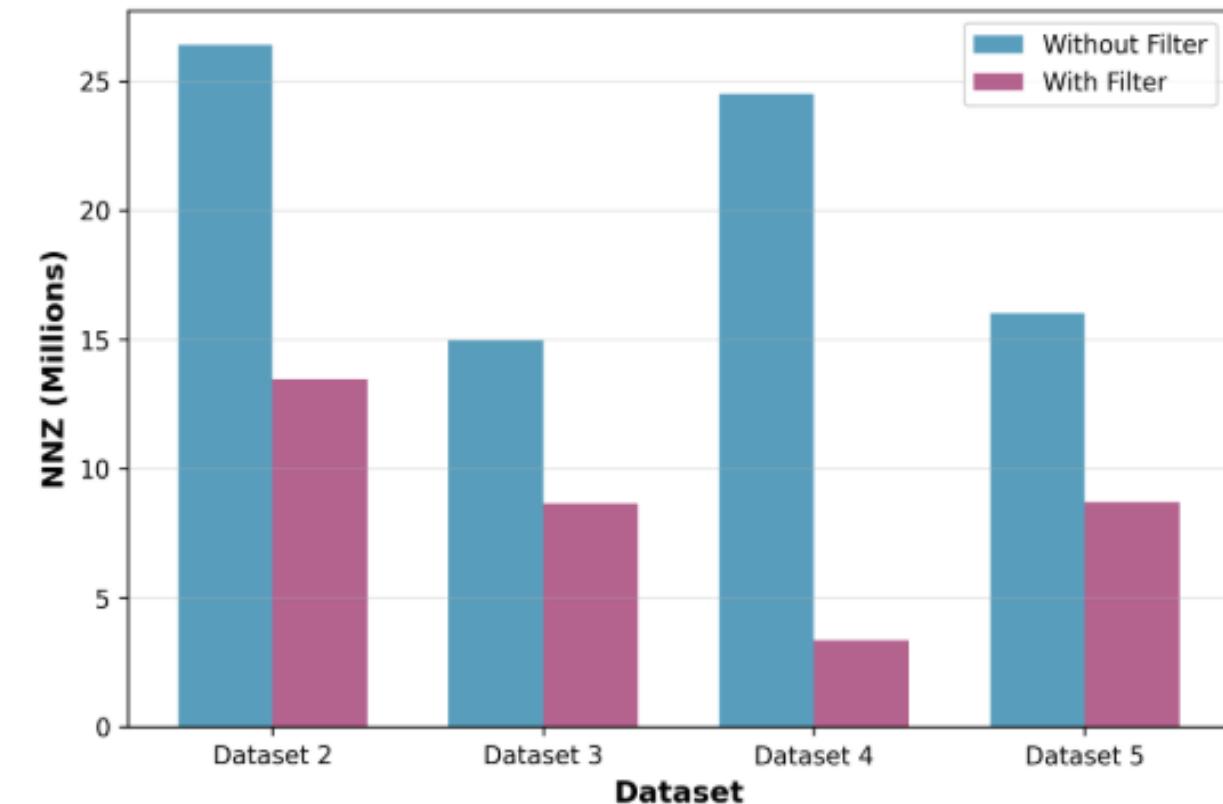
Metric	Baseline	tilepredpermspmm	Change
Compute time	53.6 ms	157.8 ms	2.9x slower
Input nnz	16.0 M	8.7 M	46% reduction
Input rows	36,706	3,660	90% reduction
Tiles	-	2,726	-
Dense tiles	-	2,726 (100%)	-
Sparse tiles	-	0 (0%)	-
Matrix density	-	0.799	-

GPU USAGE

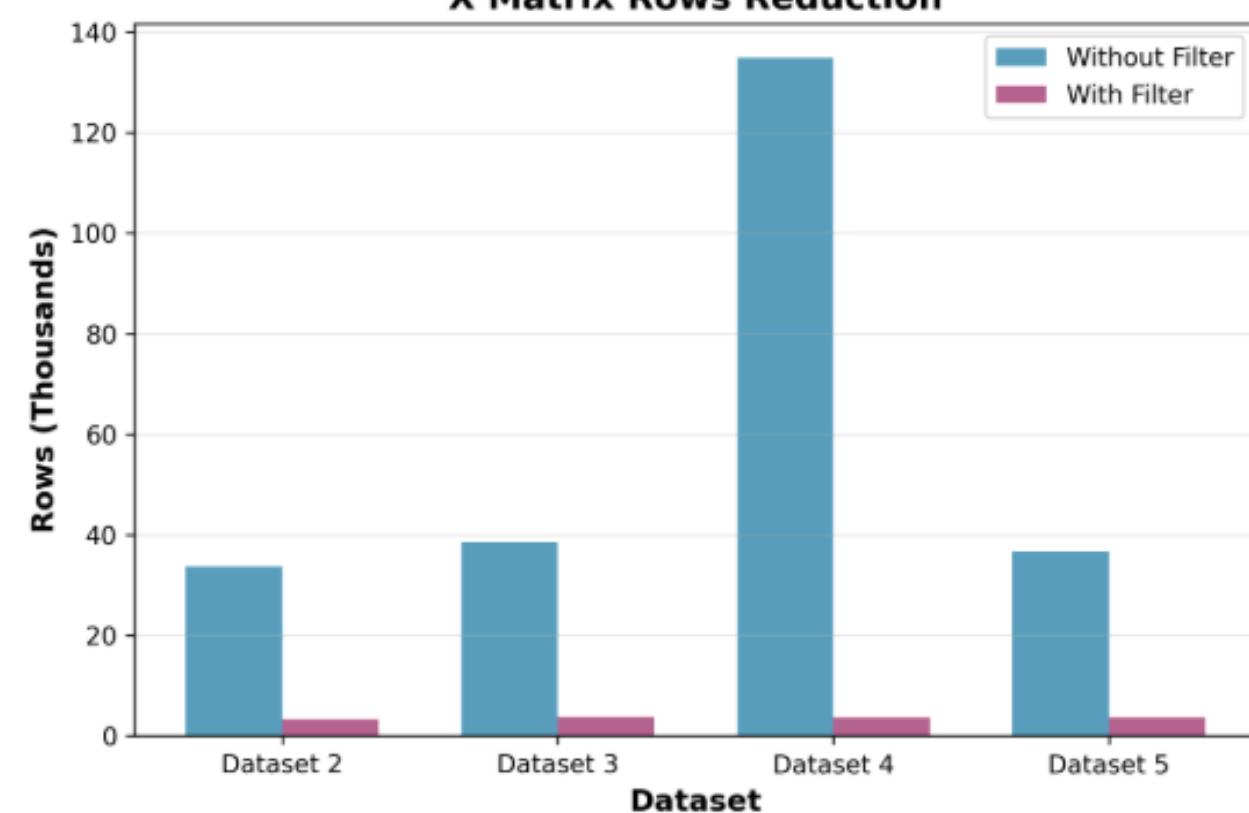
Dataset	Dense Tiles	GPU Workload	GPU Benefit
5	2,726 (100%)	High	Limited by launch overhead
2	1,049 (13.6%)	Moderate	Coordination overhead
3	105 (1.6%)	Low	Minimal
4	43 (1.7%)	Low	Minimal

X matrix Cells × Genes

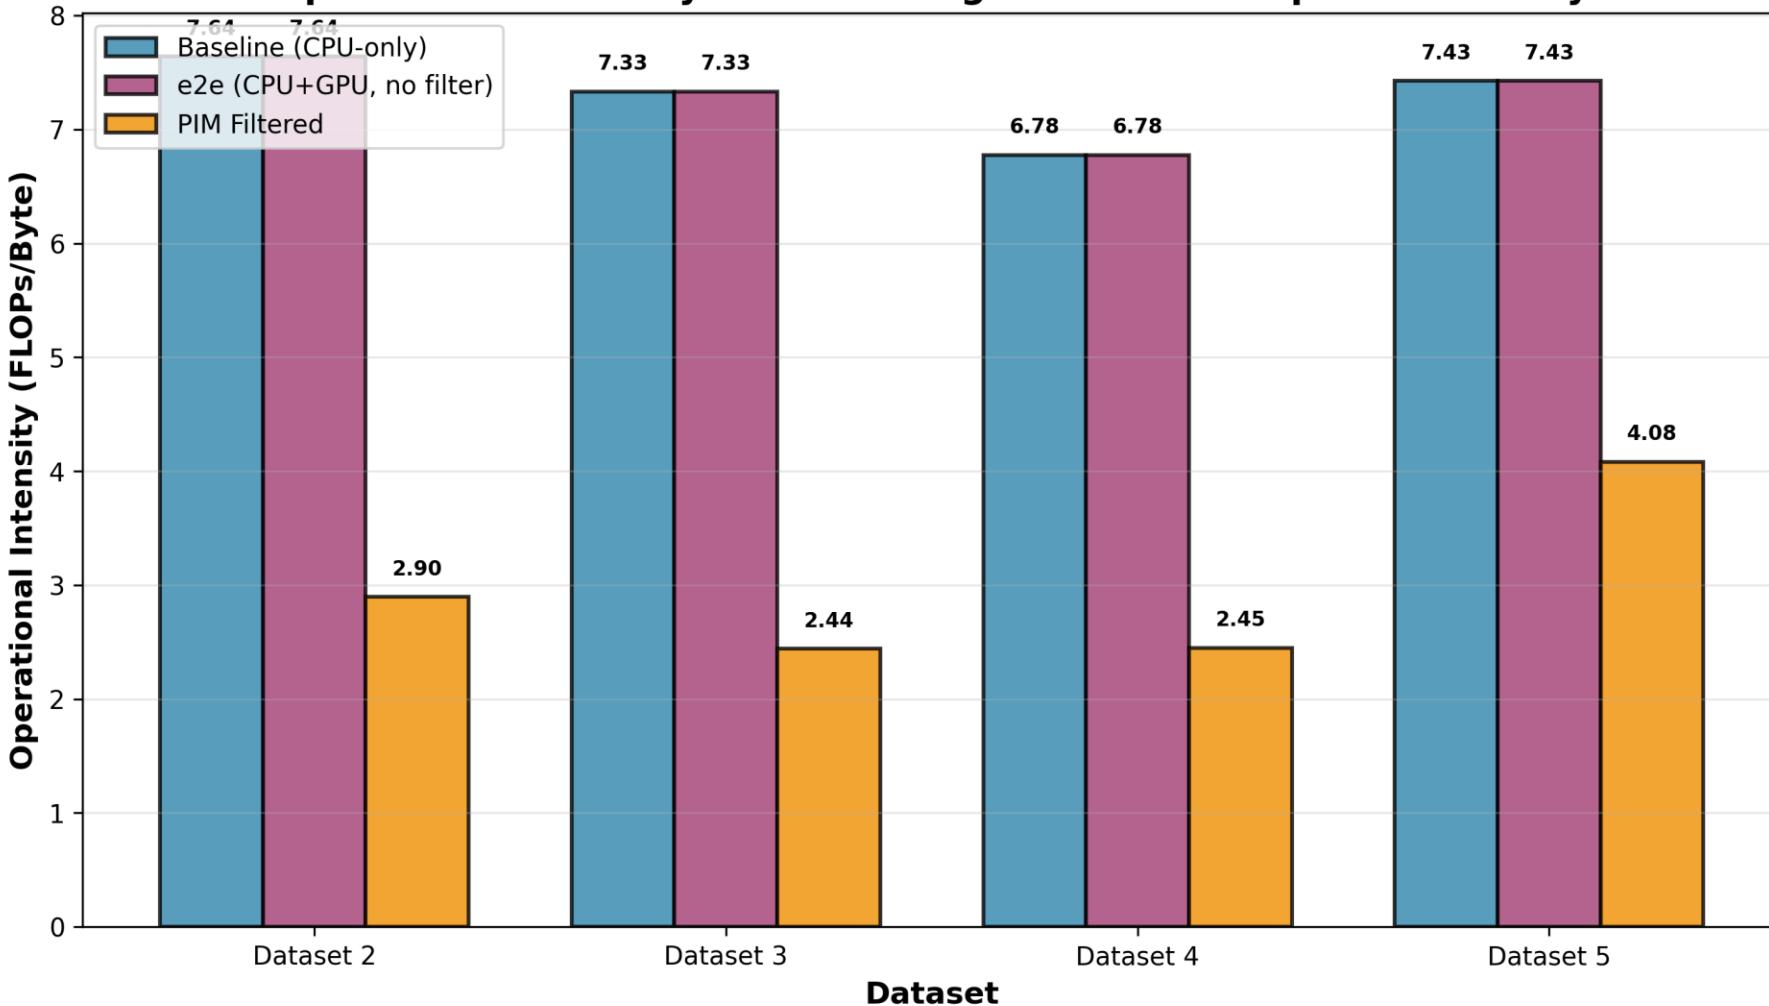
X Matrix NNZ Reduction



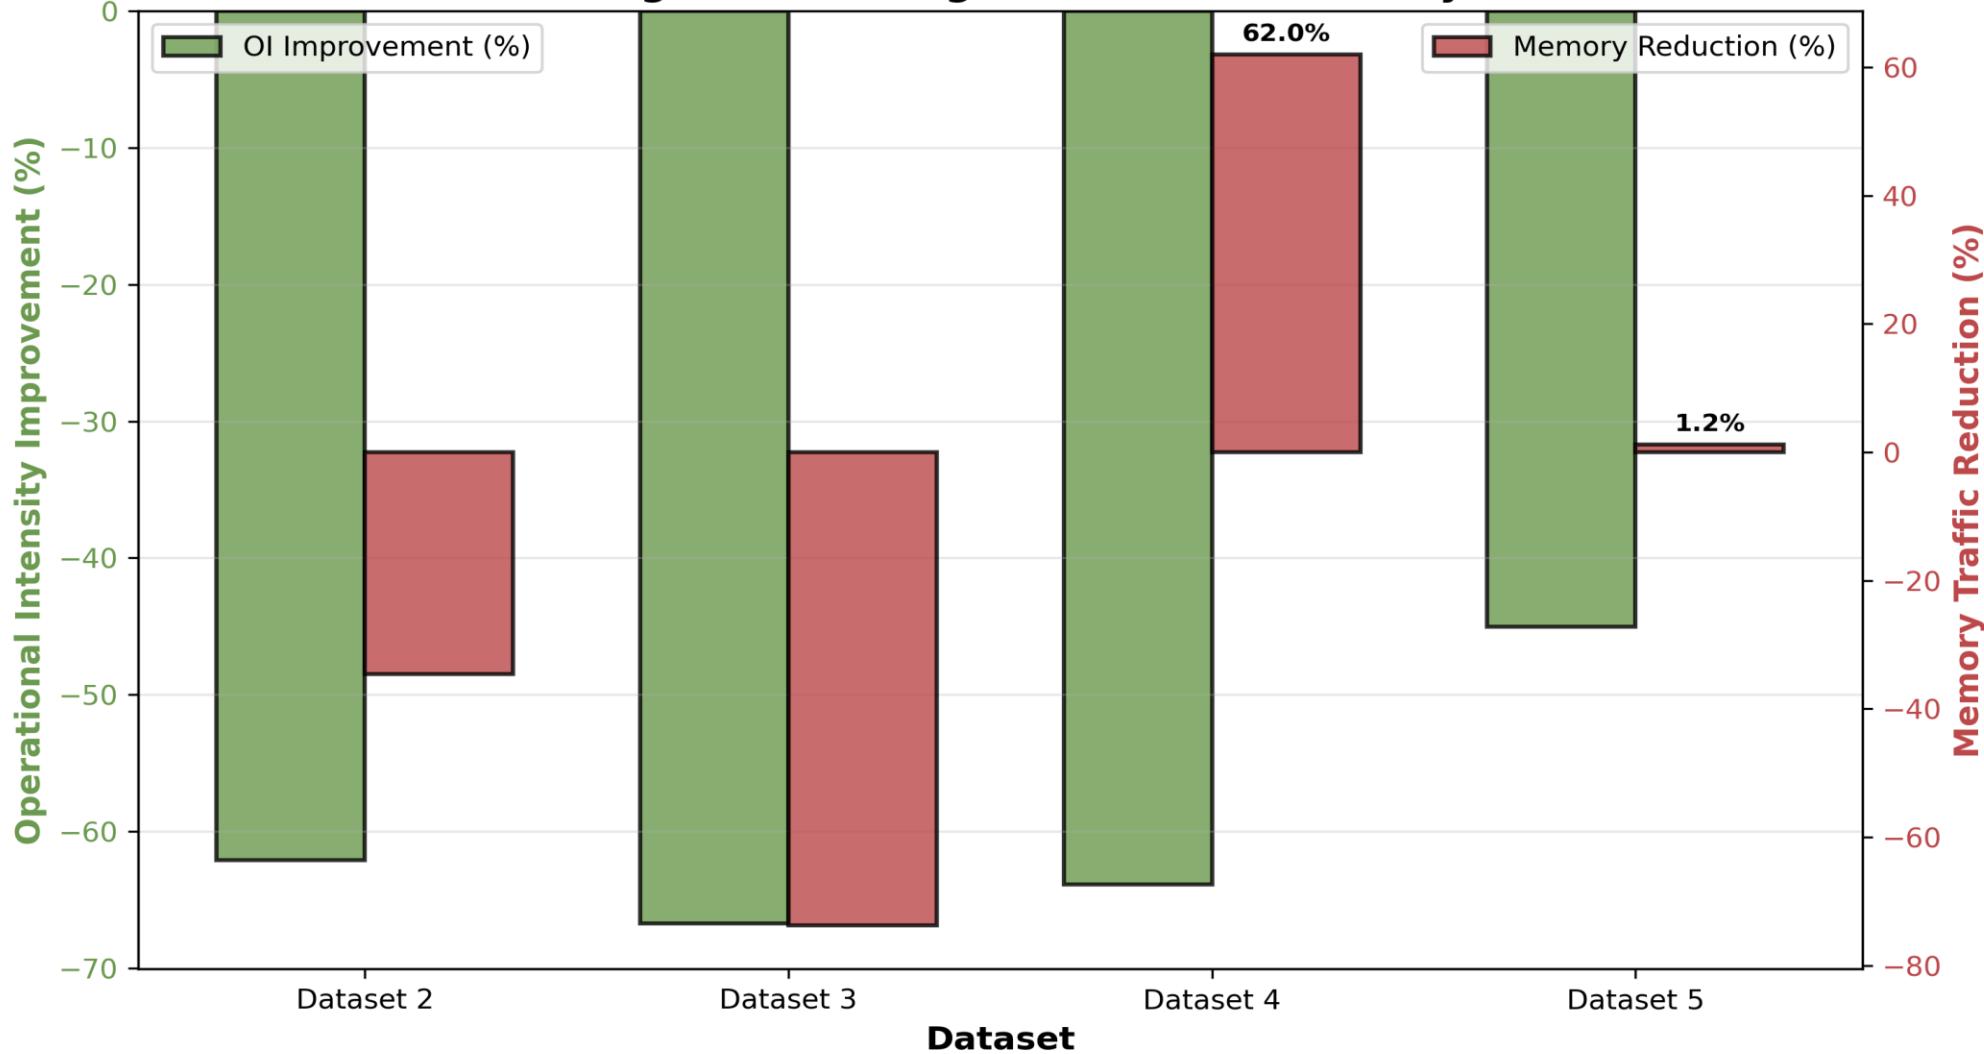
X Matrix Rows Reduction



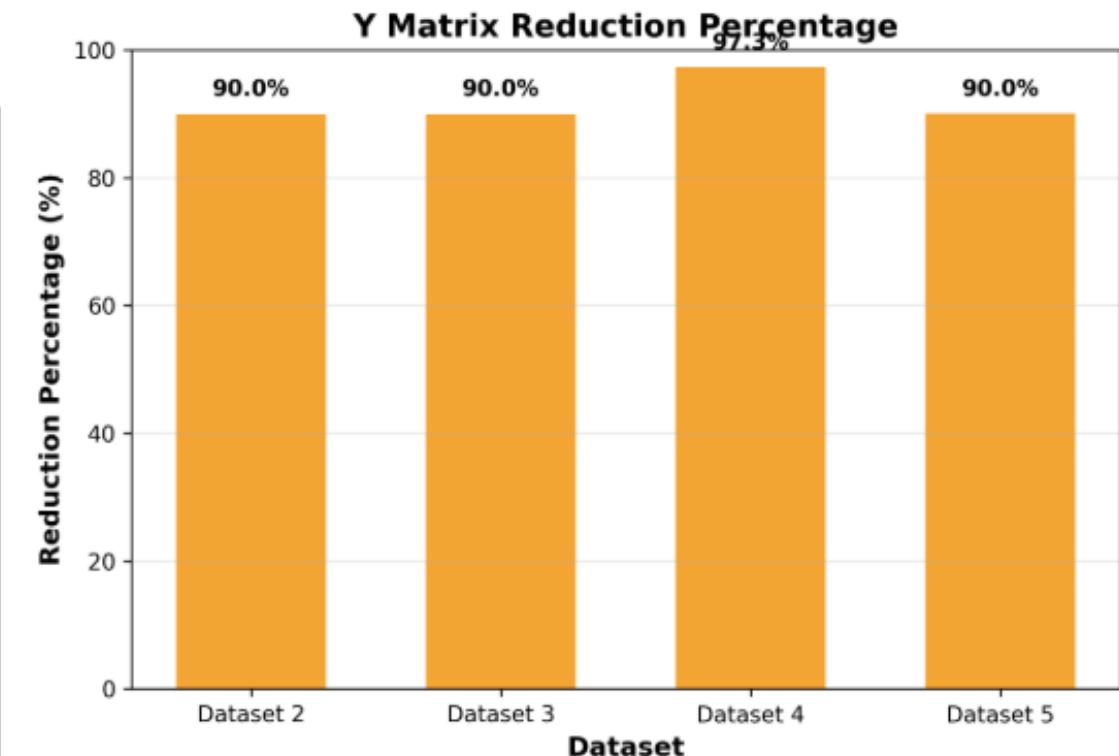
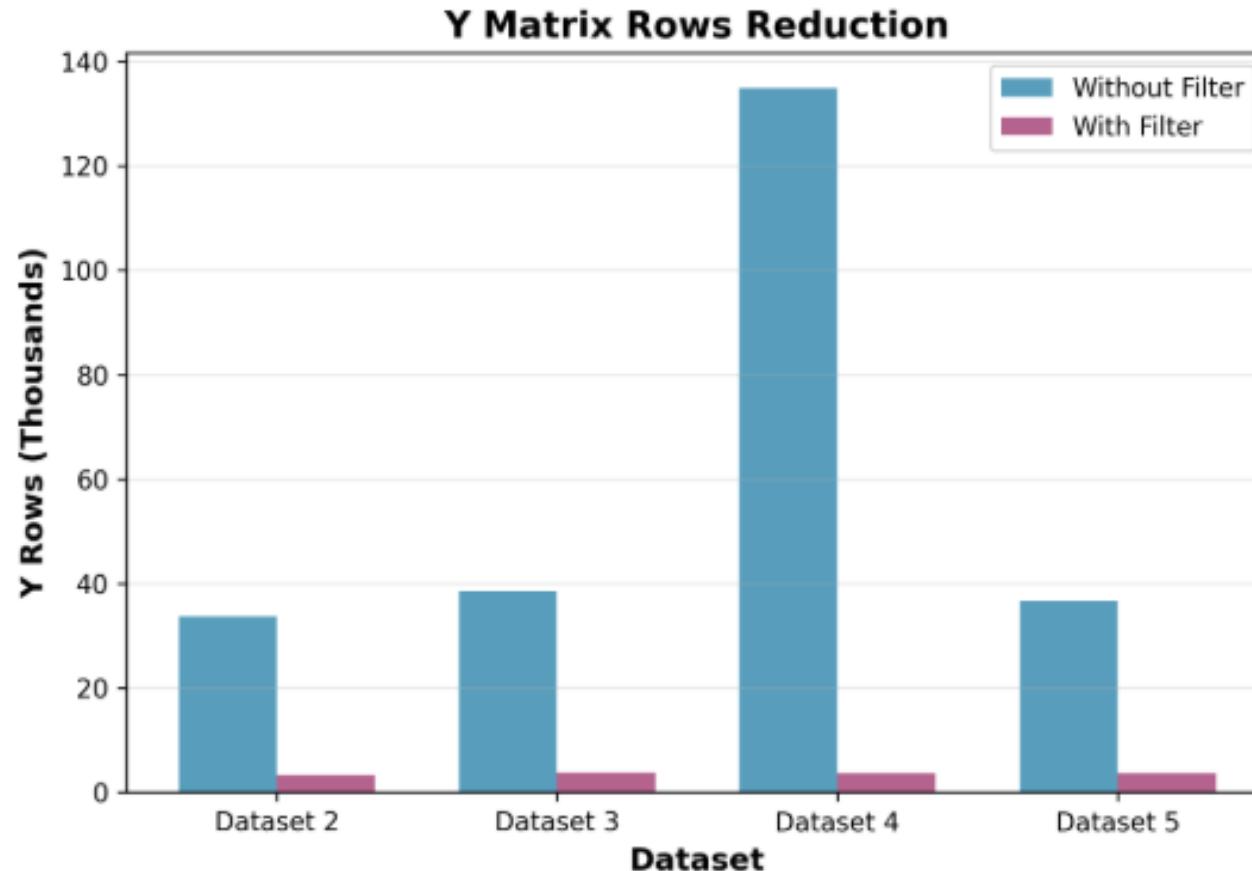
Operational Intensity: PIM Filtering Increases Compute Efficiency

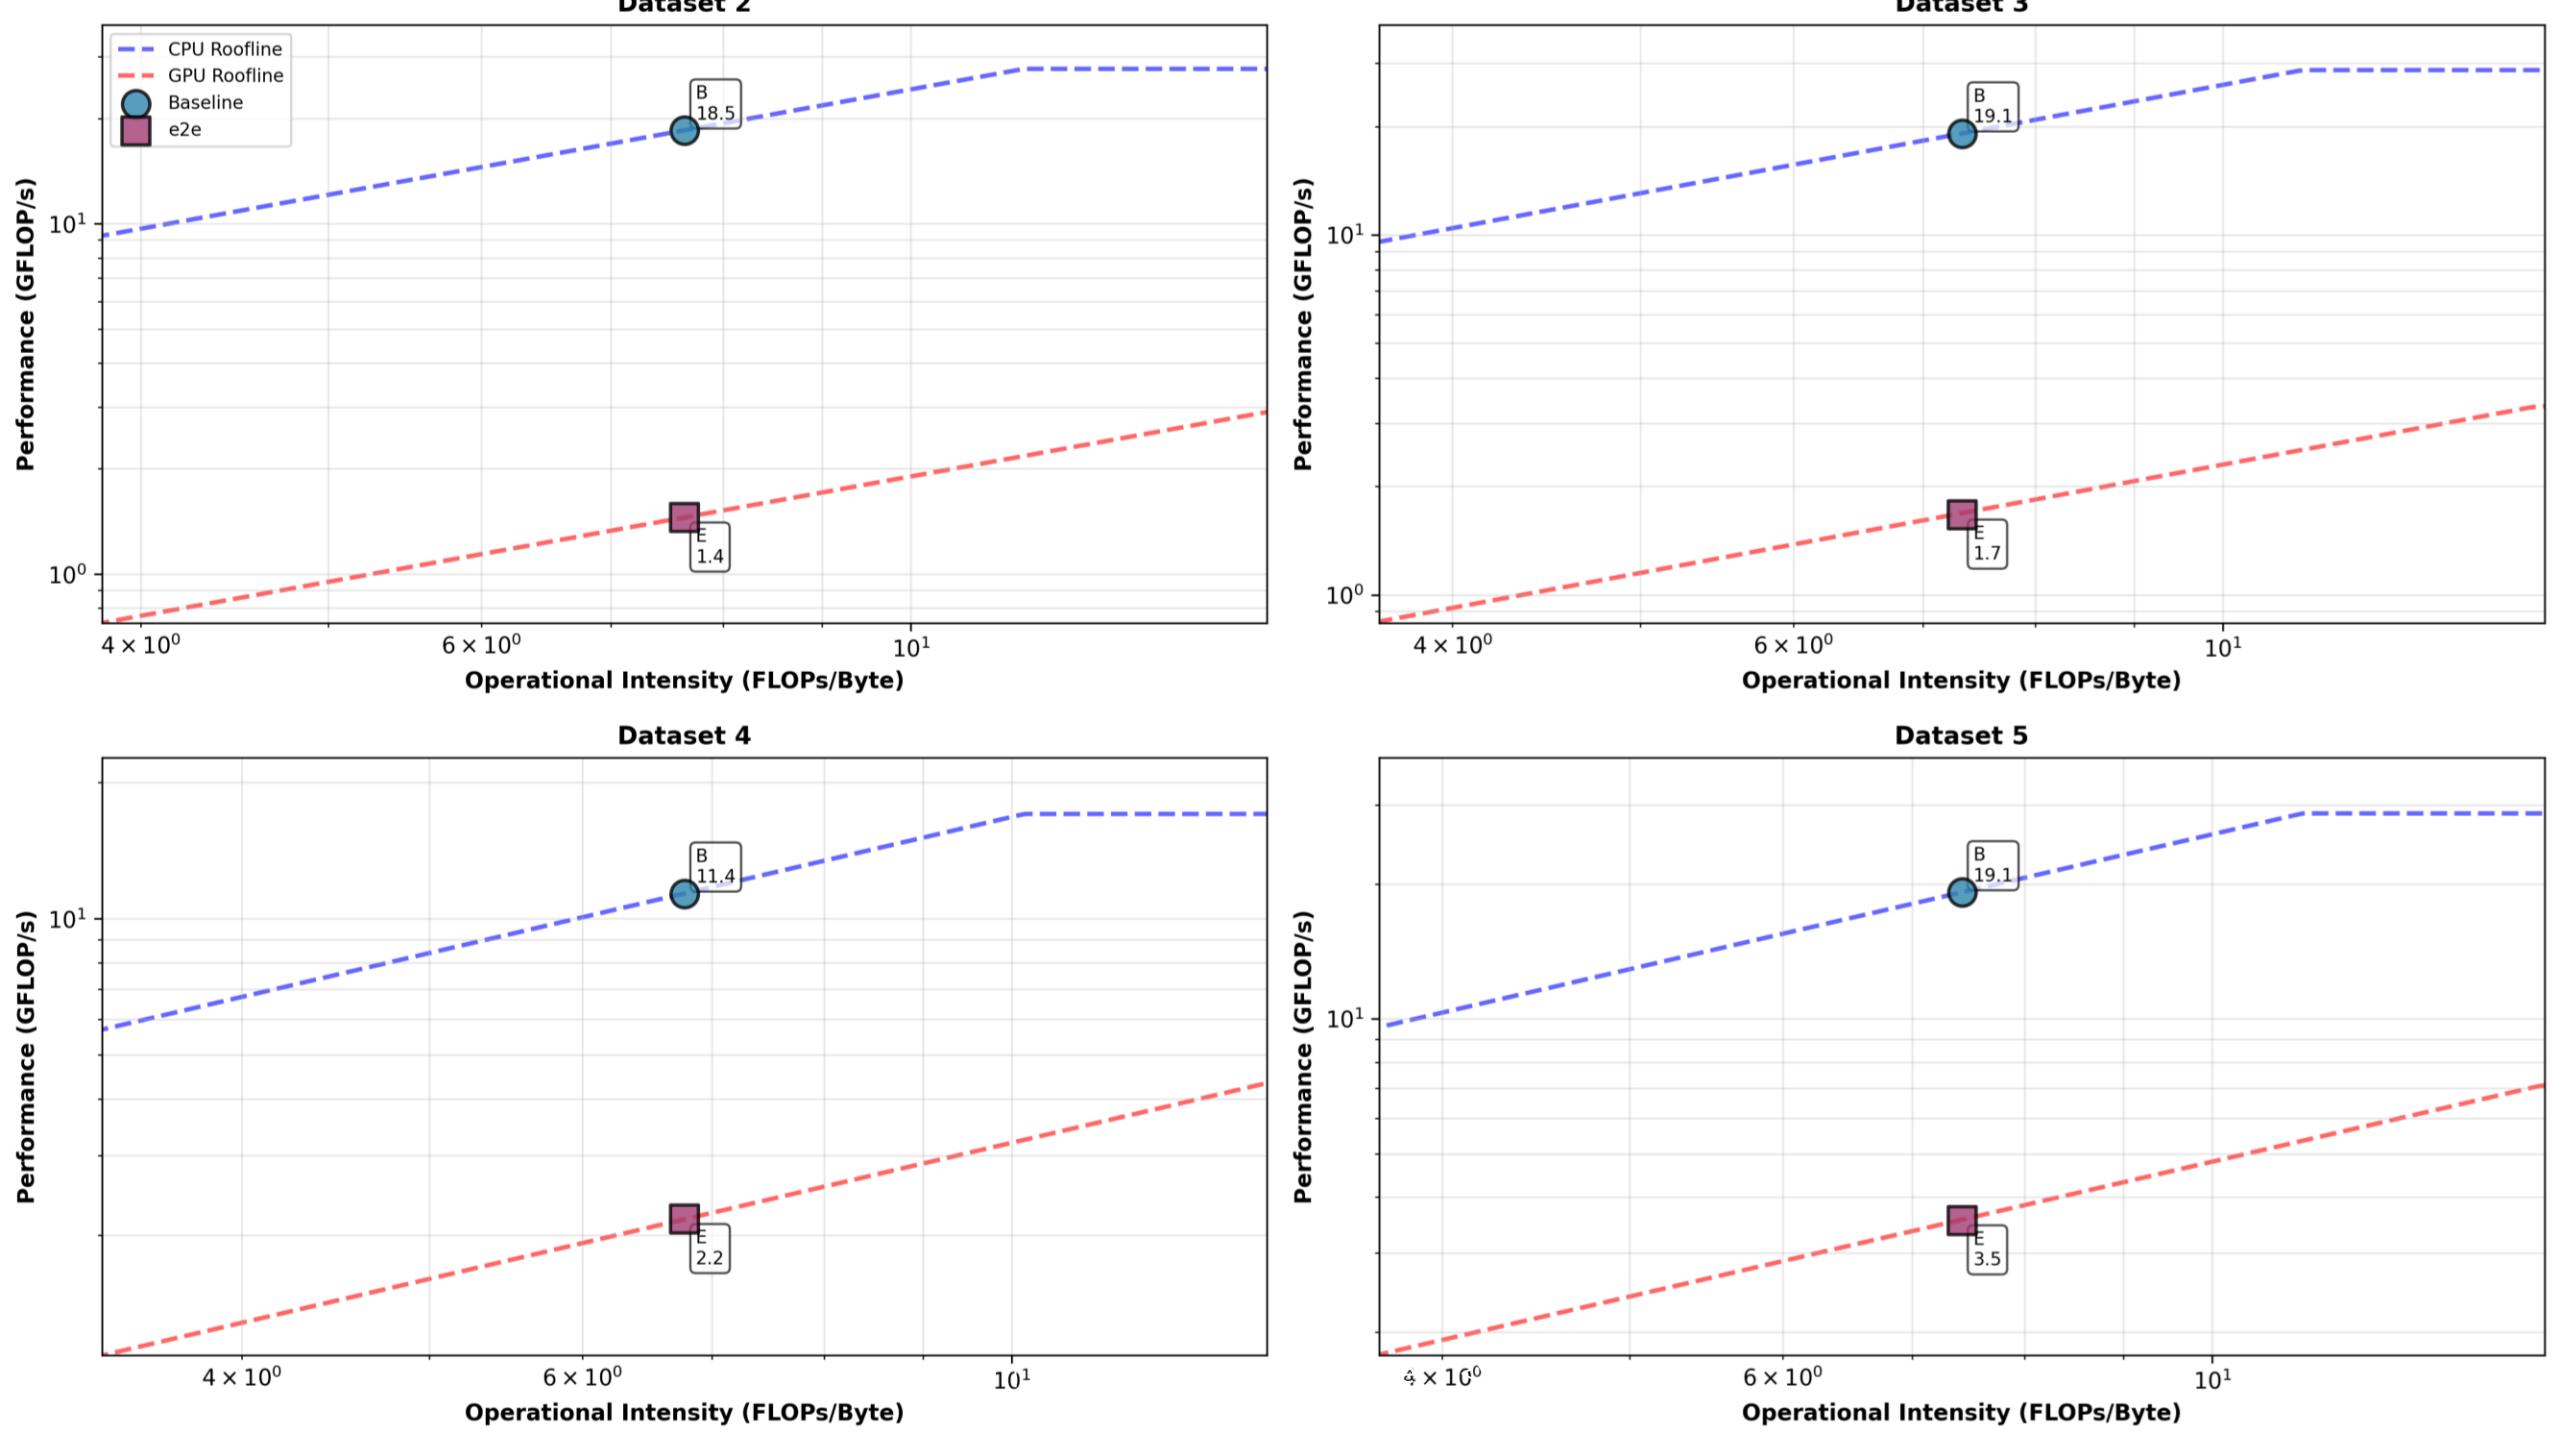


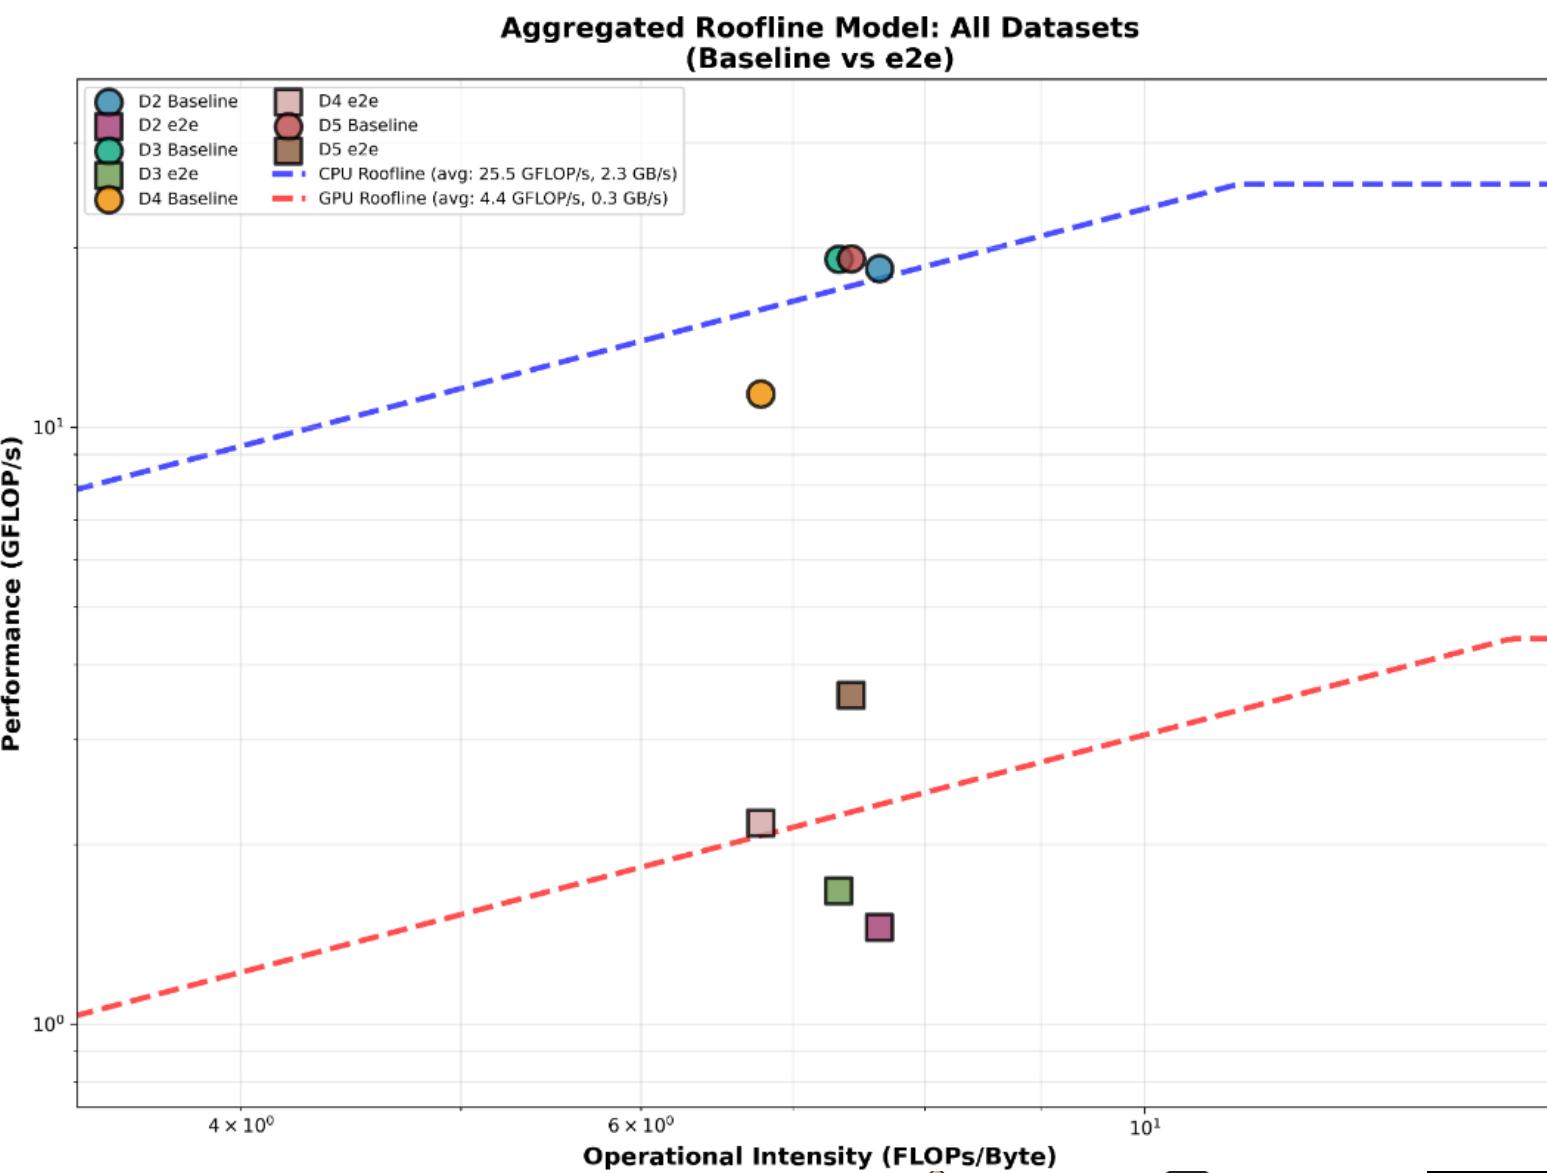
PIM Filtering Benefits: Higher OI + Less Memory Traffic



OUTPUT MATRIX $Y = X * W$







Contributions

- Demonstrates PIM filtering effectiveness with 42-86% nnz reduction across datasets.
- Provides a performance characterization.
- Reveals hybrid execution trade-offs between CPU and GPU workloads and coordination costs.

Future Work

- Adaptive tile sizing to use:
 - larger tiles for dense regions
 - smaller ones for sparse regions.
- Quantization of data.
- Look into structured sparsity
- GPU batching to aggregate multiple dense tiles into single GPU calls.
- Dynamic threshold tuning to adjust dense/sparse classification per dataset.
- Detailed overhead profiling to separately measure tiling, permutation, and transfer costs.

THANK YOU!!!

